

181840356-周娴静-实验4报告

一、题目

1. 分别编写MapReduce程序和Spark程序统计双十一最热门的商品和最受年轻人(age<30)关注的商家（“添加购物车+购买+添加收藏夹”前100名）；
2. 编写Spark程序统计双十一购买了商品的男女比例，以及购买了商品的买家年龄段的比例；
3. 基于Hive或者Spark SQL查询双十一购买了商品的男女比例，以及购买了商品的买家年龄段的比例；
4. 预测给定的商家中，哪些新消费者在未来会成为忠实客户，即需要预测这些新消费者在6个月内再次购买的概率。基于Spark MLlib编写程序预测回头客，评估实验结果的准确率。

| | |
|--------|-------------------|
| 用户行为日志 | user_log_format1 |
| 用户画像 | user_info_format1 |
| 训练数据 | train_format1 |
| 测试数据 | test_format1 |

二、设计思路及程序关键代码

1. Task1

(1) Task11 (Mapreduce)

- 类比作业5 wordcount shakespeare
- 【限定时间为双十一】首先筛选出time_stamp=1111的商品
- 【对商品计数】筛选action_type !=0的数据，以商品item_id作为它的key

```
1 public void map(Object key, Text value, Context context
2 ) throws IOException, InterruptedException {
3     StringTokenizer itr = new StringTokenizer(value.toString());
4     while (itr.hasMoreTokens()) {
5         word.set(itr.nextToken());
6         //把每一行拆解成一个一个属性,下标从0开始
7         String[] temp=word.toString().split(",");
8         //首先筛选出time_stamp=1111的商品 [5]
9         //对action_type, 如果是1, 2, 3 (不等于0) 则计数 [6]
10        if(temp[5].equals("1111"))
11            if(!temp[6].equals("0"))
```

```

12         {
13             //以商品id    item_id作为它的key    [1]
14             context.write(new Text(temp[1]), one);
15         }
16     }
17 }
18 }
19 }

```

(2) Task12 (Mapreduce)

- 与Task11类似
- 仿照作业5 (wordcount shakespeare) 设置停词文件user_info_skip.csv
- 存储patternsToSkip的HashSet, 供后续使用

```

1 public void setup(Context context) throws IOException,
  InterruptedException {
2     conf = context.getConfiguration( );
3     //caseSensitive = conf.getBoolean ( "wordcount.case.sensitive",true);
4     Path patternsPath=new
  Path("hdfs://localhost:9000/Task12/user_info_skip.csv");
5     String patternsFileName = patternsPath.getName( ).toString();
6     parseSkipFile(patternsFileName);
7     //    }
8     //}
9 }

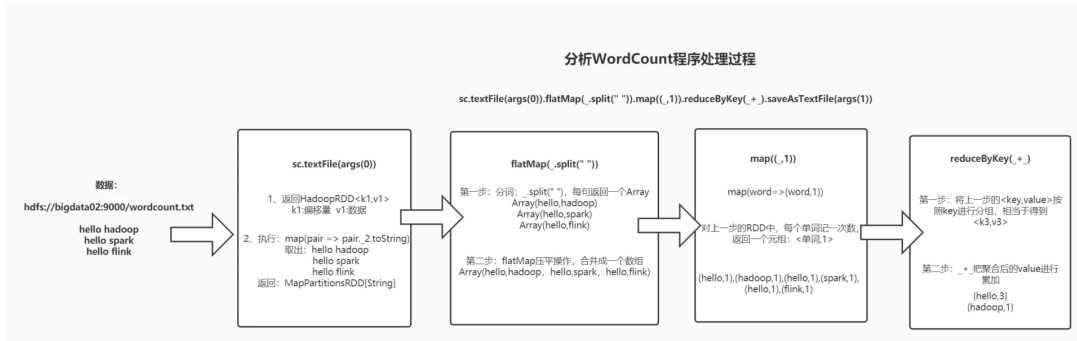
```

```

1 while(itr.hasMoreTokens()){
2     word.set(itr.nextToken());
3     //把每一行拆解成一个一个属性,下标从0开始
4     String[] temp=word.toString().split(",");
5     //首先筛选出time_stamp=1111的商品    [5]
6     //对action_type, 如果是1, 2, 3 (不等于0) 则计数    [6]
7     //购物者是年轻人    [0]
8     if(temp[5].equals("1111"))
9         if(!temp[6].equals("0"))
10             if(! patternsToSkip.contains(temp[0])){
11                 //以商家id    merchant_id作为它的key    [3]
12                 context.write(new Text(temp[3]),one);    //用
  context.write收集<key,value>对
13                 Counter counter =
  context.getCounter(CountersEnum.class.getName(),CountersEnum.INPUT_WORDS.t
  oString()); //2.0
14                 counter.increment(1);    //2.0
15             }

```

(3) Task11 (Spark)



- 读取数据文件，去除第一行
- 首先筛选出双十一（“1111”）的数据
- 进一步筛选action_type !=0的数据
- 实现二次排序

```
1 val
  data=sc.textFile("hdfs://localhost:9000/Task12/input2/user_log_format1.csv")
    .flatMap(_.split("\n"))
2 val arr = data.take(1)
5 val data1 = data.filter(!arr.contains(_)).filter(line=>line.split(",")
  (5).equals("1111")).map{
6     line => (line.split(",")(1), line.split(",")(6))
7     }.mapValues(_.toInt)
8 val data2 = data1.filter(value => value._2>0)
10 val data3 = data2.countByKey().toSeq.sortWith(_. _2>_. _2).take(100)
12 val data4 = sc.parallelize(data3)
15 data4.saveAsTextFile("hdfs://localhost:9000/Task12/output1-1")
16 :quit
18 hdfs dfs -cat /Task12/output1-1/*
```

(4) Task12 (Spark)

```
1 val log =
  spark.read.format("csv").option("header","true").load("hdfs://localhost:90
    00/Task12/input2/user_log_format1.csv")
2 val info =
  spark.read.format("csv").option("header","true").load("hdfs://localhost:90
    00/Task12/user_info_format1.csv")
5 val info1 = info.filter("age_range<4 and
  age_range>0").select("user_id","age_range")
6 val log1 = log.filter("time_stamp=1111 and
  action_type!=0").select("user_id","seller_id","action_type")
8 val dfjoin = info1.join(log1,"user_id")
10 val dfss = dfjoin.groupBy("seller_id").count()
12 //orderby count in desc and take first 100
```

```

13 val rddss = dfss.orderBy(dfss("count").desc).rdd.map(x=>
    (x(0),x(1))).take(100)
14 //save as testFile
15 sc.parallelize(rddss).saveAsTextFile("hdfs://localhost:9000/Task12/output1
    -2")
16 :quit
18 hdfs dfs -cat /Task12/outp

```

2. Task2

(1) Task21

- 基本数据筛选，筛选出双十一购买商品的数据条以及gender为（0，1）的数据
- 进行dataframe的join
- 对gender作进一步分组，并统计

```

1 val log
  =spark.read.format("csv").option("header","true").load("hdfs://localhost:9
    000/Task12/input2/user_log_format1.csv")
2 val info
  =spark.read.format("csv").option("header","true").load("hdfs://localhost:9
    000/Task12/user_info_format1.csv")
3 val log1 =
  log.select("user_id","time_stamp","action_type").filter("time_stamp=1111
    and action_type=2")
4 val info1 = info.select("user_id","gender").filter("gender=1 or gender=0")
5 val dfjoin1 = info1.join(log1,"user_id")
6 val dfcount1 = dfjoin1.groupBy("gender").count()
7 dfcount1.withColumn("ratio",dfcount1("count")/dfjoin1.count).show
8 val rddr1 =
  dfcount1.withColumn("ratio",dfcount1("count")/dfjoin1.count).select("gende
    r","ratio").rdd.map(x=>(x(0),x(1))).collect()
9 sc.parallelize(rddr1).saveAsTextFile("hdfs://localhost:9000/Task12/output2
    -1")

```

(2) Task22

- 与Task21类似

```

1 val log
  =spark.read.format("csv").option("header","true").load("hdfs://localhost:9
    000/Task12/input2/user_log_format1.csv")
2 val info
  =spark.read.format("csv").option("header","true").load("hdfs://localhost:9
    000/Task12/user_info_format1.csv")

```

```

5 val info1 = info.select("user_id","age_range").filter("age_range>0 and
age_range<9")
6 val log1 =
log.select("user_id","time_stamp","action_type").filter("time_stamp=1111
and action_type=2")
8 val dfjoin = info1.join(log1,"user_id")
10 val dfcount = dfjoin.groupBy("age_range").count()
12 dfcount.withColumn("ratio",dfcount("count")/dfjoin.count()).show
15 /*将答案保存到本地*/
16 val dfratio =
dfcount.withColumn("ratio",dfcount("count")/dfjoin.count()).select("age_rang
e","ratio").orderBy("age_range")
18 val rddr = dfratio.rdd.map(x=>(x(0),x(1)))
20 val rddr = dfratio.rdd.map(x=>(x(0),x(1))).collect()
22 sc.parallelize(rddr).saveAsTextFile("hdfs://localhost:9000/Task12/output2-
2")

```

3. Task3

(1) Task31

- 将数据存储为dataframe以后进行createOrReplaceTempView试图转化，方便后续进行表的连接
- 采用嵌套查询方法，内部select初步筛选出购买了商品且gender为（0，1）的数据，设置distinct保证最多只读取该用户id一次
- 外部select进一步根据gender分组，并进行统计

```

1 //terminal
2 cd $SPARK_HOME
3 bin/spark-shell
4
5
6 //scala
7 val user_log =
spark.read.format("csv").option("header","true").load("hdfs://localhost:90
00/Task12/input2/user_log_format1.csv")
8
9 val user_info = spark.read.format("csv").option("header","true").load("hdf
s://localhost:9000/Task12/user_info_format1.csv")
10
11 user_log.createOrReplaceTempView("table_log")
12
13 user_info.createOrReplaceTempView("table_info")
15 //3-1
17 val people1=spark.sql("select gender,count(*) as num from (select distinct
a.user_id,gender from table_log a,table_info b where a.user_id=b.user_id
and a.action_type='2' and gender in('0','1')) group by gender")

```

```

18 people1.show
19 //下面显示ratio
20 people1.agg("num"-> "sum").show //结果为407308
21 people1.withColumn("ratio",people1("num")/407308).show

```

(2) Task32

- 与Task31类似

```

1 //3-2
2 val people2=spark.sql("select age_range,count(*) as num from (select
distinct a.user_id,age_range from table_log a,table_info b
where a.user_id=b.user_id and a.action_type='2'
and age_range in('1','2','3','4','5','6','7','8')) group
by age_range order by age_range")
3 people2.show
4 // //下面显示ratio(但不是sql语句
5 people2.agg("num"-> "sum").show //结果为329039
6 people2.withColumn("ratio",people2("num")/329039).show

```

4. Task4

参考：

<https://www.jianshu.com/p/aa6cb1ef6f69>

<https://tianchi.aliyun.com/notebook-ai/detail?postId=143593>

- 参考Baseline做法，选取age_range、gender、total_logs、unique_item_ids、categories、browse_days、one_clicks、shopping_carts、purchase_times、favourite_times 建立特征工程
- 对建立好的特征作缺失值处理
- 利用pyspark-ml随机森林进行预测

```

1 #jupyter notebook
2 # In[1]:
3 from pyspark.sql import SparkSession
4 spark = SparkSession.builder.appName('random_forest').getOrCreate()
5 df=spark.read.csv('hdfs://localhost:9000/Task12/feature.csv',inferSchema=True,header=True)
6 # In[2]:
7 from pyspark.ml.feature import VectorAssembler
8 df_assembler = VectorAssembler(inputCols=
9 ['age_range','gender','total_logs','unique_item_ids','categories','browse_
10 days','one_clicks','shopping_carts','purchase_times','favourite_times'], o
11 utputCol="features")
12 df = df_assembler.transform(df)

```

```

15 -测试集和训练集划分比例: 0.75: 0.25
16 -随机森林基学习器数量: 120
17 # In[3]:
18 model_df = df.select(['features','label'])
19 train_df,test_df=model_df.randomSplit([0.75,0.25])
20 # In[4]:
22 from pyspark.ml.classification import RandomForestClassifier
23 rf_classifier=RandomForestClassifier(labelCol='label',numTrees=120).fit(tr
ain_df)
24 rf_predictions=rf_classifier.transform(test_df)
26 # In[5]:
27 from pyspark.ml.evaluation import BinaryClassificationEvaluator
28 from pyspark.ml.evaluation import MulticlassClassificationEvaluator
29 rf_accuracy=MulticlassClassificationEvaluator(labelCol='label',metricName=
'accuracy').evaluate(rf_predictions)
31 print('MulticlassClassificationEvaluator 随机森林测试的准确
性: {0:.0%}'.format(rf_accuracy))
32 # In[6]:
34 rf_auc=BinaryClassificationEvaluator(labelCol='label').evaluate(rf_predict
ions)
35 print('BinaryClassificationEvaluator 随机森林测试的准确
性: {0:.0%}'.format(rf_auc))
36

```

三、Ubuntu命令行关键代码

1. Mapreduce指令

```

1 cd ~/hadoop_installs/hadoop-3.3.0
2 sbin/start-all.sh
3 jps
5 cd /mnt/hgfs/ubuntu_share/Experiment4/No1.2
6 bin/hdfs dfs -mkdir /Task12
8 #input1放的是mytest.CSV测试样例
9 bin/hdfs dfs -mkdir /Task12/input1
10 bin/hdfs dfs -put
   /mnt/hgfs/ubuntu_share/Experiment4/No1.2/user_info_skip.csv /Task12/
12 bin/hdfs dfs -put /mnt/hgfs/ubuntu_share/Experiment4/No1.2/mytest11.CSV
   /Task12/input1
15 bin/hadoop jar /mnt/hgfs/ubuntu_share/Experiment4/No1.2/popular_merchant-
7.0-SNAPSHOT.jar project.code /Task12/input2 /Task12/output7

```

2. 安装Spark

```

1 tar zxvf spark-3.0.1-bin-hadoop2.7.tgz -C /usr/local/
2 cd /usr/local/

```

```
3 mv spark-3.0.1-bin-hadoop2.7 spark
5 cd /mnt/hgfs/ubuntu_share/
6 tar zxvf scala-2.12.11.tgz -C /usr/local/
7 cd /usr/local/
8 mv scala-2.12.11 scala
9 vim ~/.bashrc
10 # ~/.bashrc内容
12 #SCALA
13 export SCALA_HOME=/usr/local/scala
14 export PATH=$PATH:$HOME/bin:$JAVA_HOME/bin:$SCALA_HOME/bin
15 #SPARK
16 export SPARK_HOME=/usr/local/spark
17 export PATH=$PATH:$SPARK_HOME/bin
18 #localhost:8080
20 cd /usr/local/spark/conf
22 cp spark-env.sh.template spark-env.sh
23 vim spark-env.sh
25 #内容
26 export JAVA_HOME=/usr/java/jdk1.8.0_261
27 export HADOOP_HOME=/usr/local/hadoop
28 export HADOOP_CONF_DIR=/usr/local/hadoop/etc/hadoop
29 export SCALA_HOME=/usr/local/scala
30 export SPARK_HOME=/usr/local/spark
31 export SPARK_MASTER_IP=127.0.0.1
32 export SPARK_MASTER_PORT=7077
33 export SPARK_MASTER_WEBUI_PORT=8099
34 export SPARK_WORKER_CORES=3
35 export SPARK_WORKER_INSTANCES=1
36 export SPARK_WORKER_MEMORY=5G
37 export SPARK_WORKER_WEBUI_PORT=8081
38 export SPARK_EXECUTOR_CORES=1
39 export SPARK_EXECUTOR_MEMORY=1G
40 export LD_LIBRARY_PATH=${LD_LIBRARY_PATH}:$HADOOP_HOME/lib/native
```

四、实验结果及截图

1. Task1截图

Task11截图


```
root@zxj13-virtual-machine:/usr/local/spark#hdfs dfs -cat /Task12/output1-1/*
(191499,2494)
(353560,2250)
(1059899,1917)
(713695,1754)
(655904,1674)
(67897,1572)
(221663,1547)
(1039919,1511)
(454937,1387)
(81360,1361)
(514725,1356)
(783997,1351)
(823766,1343)
(107407,1319)
(889095,1272)
(936203,1270)
(770668,1257)
(698879,1235)
(349999,1218)
(671759,1167)
(186456,1162)
(315345,1067)
(729259,1021)
(946001,1015)
(181387,1002)
(926069,1002)
(28895,983)
(89953,975)
(413046,965)
(944554,948)
(617878,927)
(676215,873)
(213297,864)
(15207,859)
(513855,842)
(49881,842)
(48664,831)
(179830,827)
(981145,815)
(441588,814)
(939915,805)
```

| | |
|---------------|--|
| (764906,787) | |
| (1025501,778) | |
| (3001,765) | |
| (147751,763) | |
| (343432,758) | |
| (141675,757) | |
| (100215,745) | |
| (722301,742) | |
| (678194,732) | |
| (526229,730) | |
| (1100222,730) | |
| (952198,725) | |
| (846908,711) | |
| (1102377,708) | |
| (605027,707) | |
| (201405,704) | |
| (15173,698) | |
| (758374,698) | |
| (335720,695) | |
| (110347,693) | |
| (950987,687) | |
| (834362,687) | |
| (399879,685) | |
| (81901,676) | |
| (487805,668) | |
| (203050,659) | |
| (1075577,659) | |
| (173776,659) | |
| (796566,653) | |
| (276750,650) | |
| (209821,647) | |
| (735931,647) | |
| (779070,645) | |
| (235204,640) | |
| (318890,635) | |
| (986262,634) | |
| (886674,622) | |
| (386646,616) | |
| (717309,616) | |
| (28186,615) | |
| (376482,610) | |
| (554408,603) | |
| (772645,601) | |
| (992011,598) | |
| (784134,597) | |
| (472166,595) | |
| (825218,590) | |
| (566407,585) | |
| (918348,580) | |
| (982357,580) | |
| (293244,577) | |
| (419724,571) | |
| (256896,559) | |
| (893999,554) | |
| (870470,549) | |
| (1042707,549) | |
| (82431,548) | |
| (1093758,545) | |
| (1112049,543) | |

Task12截图

```
root@zxj13-virtual-machine:/usr/local/spark#hdfs dfs -cat /Task12/output1-2/*
```

```
(4044,7278)
(3491,3661)
(1102,3588)
(3828,3434)
(4173,3348)
(3734,3303)
(2385,3214)
(4976,3064)
(798,2997)
(422,2893)
(1892,2792)
(1393,2774)
(4282,2737)
(1535,2720)
(4760,2669)
(4644,2607)
(3760,2600)
(184,2406)
(598,2400)
(3698,2111)
(375,2078)
(4043,2077)
(2537,1949)
(1760,1935)
(2482,1920)
(4659,1843)
(2138,1841)
(606,1818)
(1257,1775)
(4218,1760)
(4538,1639)
(141,1633)
(4918,1596)
(3826,1588)
(2031,1584)
(420,1579)
(2813,1553)
```

```
(173,1496)
(962,1495)
(2223,1471)
(4079,1467)
(1816,1462)
(2217,1442)
(2403,1436)
(66,1432)
(742,1423)
(2468,1360)
(1056,1356)
(4845,1344)
(2418,1339)
(2669,1335)
(4048,1333)
(4648,1263)
(1861,1253)
(4818,1220)
(4766,1203)
(2273,1198)
(3022,1194)
(2336,1186)
(361,1176)
(474,1174)
(4798,1167)
(2545,1159)
(1087,1133)
(4847,1126)
(2954,1119)
(1346,1106)
(4871,1098)
(3578,1088)
(2387,1076)
(3971,1074)
(4605,1056)
(2676,1051)
(1480,1049)
(2823,1047)
(4127,1040)
```

```
(1,1030)
(2206,1027)
(4257,1001)
(4129,969)
(4160,968)
(2193,963)
(2664,954)
(1727,951)
(310,948)
(1310,930)
(1487,918)
(1200,914)
(3859,914)
(2928,897)
(3163,885)
(3623,884)
(4287,875)
(1867,868)
(3173,867)
(2318,865)
(2677,865)
(4427,864)
(786,839)
(643,827)
```

2. Task2截图

Task21截图

```
scala> val log = spark.read.format("csv").option("header","true").load("hdfs://local
host:9000/Task12/input2/user_log_format1.csv")
log: org.apache.spark.sql.DataFrame = [user_id: string, item_id: string ... 5 more
fields]

scala> val info = spark.read.format("csv").option("header","true").load("hdfs://loca
lhost:9000/Task12/user_info_format1.csv")
info: org.apache.spark.sql.DataFrame = [user_id: string, age_range: string ... 1 mo
re field]

scala> val log1 = log.select("user_id","time_stamp","action_type").filter("time_sta
mp=1111 and action_type=2")
log1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [user_id: string, ti
me_stamp: string ... 1 more field]

scala> val info1 = info.select("user_id","gender").filter("gender=1 or gender=0")
info1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [user_id: string, g
ender: string]

scala> val dfjoin1 = info1.join(log1,"user_id")
dfjoin1: org.apache.spark.sql.DataFrame = [user_id: string, gender: string ... 2 mo
re fields]

scala> val dfcount1 = dfjoin1.groupBy("gender").count()
dfcount1: org.apache.spark.sql.DataFrame = [gender: string, count: bigint]

scala> dfcount1.withColumn("ratio",dfcount1("count")/dfjoin1.count).show
+-----+-----+-----+
|gender| count|      ratio|
+-----+-----+-----+
|    0|846054|0.7232596926427983|
|    1|323725|0.2767403073572017|
+-----+-----+-----+

scala> val rddr1 = dfcount1.withColumn("ratio",dfcount1("count")/dfjoin1.count).sel
ect("gender","ratio").rdd.map(x=>(x(0),x(1))).collect()
rddr1: Array[(Any, Any)] = Array((0,0.7232596926427983), (1,0.2767403073572017))

root@zxj13-virtual-machine: /usr/local/spark#hdfs dfs -cat /Task12/output2-1/*
(0,0.7232596926427983)
(1,0.2767403073572017)
```

Task22截图

```
scala> val log = spark.read.format("csv").option("header", "true").load("hdfs://localhost:9000/Task12/input2/user_log_format1.csv")
log: org.apache.spark.sql.DataFrame = [user_id: string, item_id: string ... 5 more fields]

scala> val info = spark.read.format("csv").option("header", "true").load("hdfs://localhost:9000/Task12/user_info_format1.csv")
info: org.apache.spark.sql.DataFrame = [user_id: string, age_range: string ... 1 more field]

scala> val info1 = info.select("user_id", "age_range").filter("age_range > 0 and age_range < 9")
info1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [user_id: string, age_range: string]

scala> val log1 = log.select("user_id", "time_stamp", "action_type").filter("time_stamp = 1111 and action_type = 2")
log1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [user_id: string, time_stamp: string ... 1 more field]

scala> val dfjoin = info1.join(log1, "user_id")
dfjoin: org.apache.spark.sql.DataFrame = [user_id: string, age_range: string ... 2 more fields]

scala> val dfcount = dfjoin.groupBy("age_range").count()
dfcount: org.apache.spark.sql.DataFrame = [age_range: string, count: bigint]

scala> dfcount.withColumn("ratio", dfcount("count")/dfjoin.count()).show
+-----+-----+-----+
|age_range|count|ratio|
+-----+-----+-----+
|7|19363|0.019736191647869567|
|3|327758|0.33407502464093547|
|8|3476|0.003542994482672861|
|5|133480|0.13605261897214427|
|6|103774|0.10577408211878409|
|1|54|5.504076584129301E-5|
|4|268549|0.2737248634428407|
|2|124637|0.12703918392891178|
+-----+-----+-----+

scala> val dfratio = dfcount.withColumn("ratio", dfcount("count")/dfjoin.count()).select("age_range", "ratio").orderBy("age_range")
dfratio: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [age_range: string, ratio: double]

scala> val rddr = dfratio.rdd.map(x=>(x(0), x(1))).collect()
rddr: Array[(Any, Any)] = Array((1, 5.504076584129301E-5), (2, 0.12703918392891178), (3, 0.33407502464093547), (4, 0.2737248634428407), (5, 0.13605261897214427), (6, 0.10577408211878409), (7, 0.019736191647869567), (8, 0.003542994482672861))

scala>

scala> sc.parallelize(rddr).saveAsTextFile("hdfs://localhost:9000/Task12/output2-2")
```

```
root@zxj13-virtual-machine: /usr/local/spark# hdfs dfs -cat /Task12/output2-2/*
(1,5.504076584129301E-5)
(2,0.12703918392891178)
(3,0.33407502464093547)
(4,0.2737248634428407)
(5,0.13605261897214427)
(6,0.10577408211878409)
(7,0.019736191647869567)
(8,0.003542994482672861)
```

3. Task3截图

Task31截图

```
scala> people1.show
+-----+-----+
|gender|   num|
+-----+-----+
|     0|285638|
|     1|121670|
+-----+-----+

scala> people1.agg("num" -> "sum").show
+-----+
|sum(num)|
+-----+
|  407308|
+-----+

scala> people1.withColumn("ratio",people1("num")/407308).show
+-----+-----+-----+
|gender|   num|          ratio|
+-----+-----+-----+
|     0|285638|0.7012825674919225|
|     1|121670|0.2987174325080774|
+-----+-----+-----+
```

Task32截图

```
scala> people2.show
+-----+-----+
|age_range|   num|
+-----+-----+
|         1|    24|
|         2|  52871|
|         3|111654|
|         4|  79991|
|         5|  40777|
|         6|  35464|
|         7|   6992|
|         8|   1266|
+-----+-----+

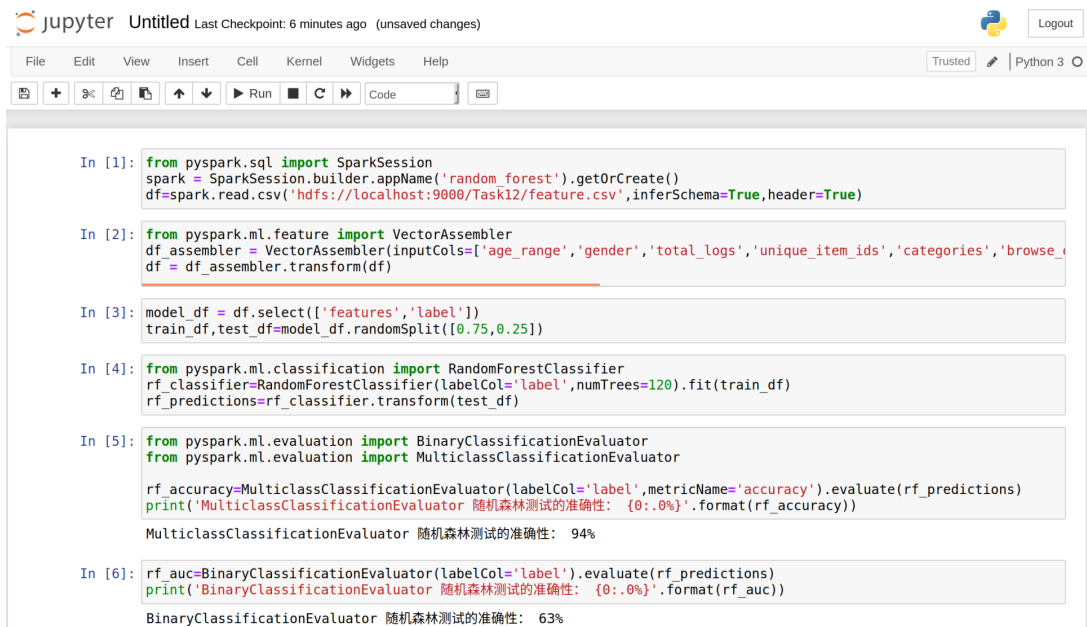
scala>

scala> people2.agg("num" -> "sum").show //结果为329039
+-----+
|sum(num)|
+-----+
|  329039|
+-----+

scala>

scala> people2.withColumn("ratio",people2("num")/329039).show
+-----+-----+-----+
|age_range|   num|          ratio|
+-----+-----+-----+
|         1|    24|7.293968192220375E-5|
|         2|  52871| 0.16068308012120144|
|         3|111654| 0.3393336352225724|
|         4|  79991| 0.24310492069329168|
|         5|  40777| 0.1239275587392376|
|         6|  35464| 0.10778053665370975|
|         7|   6992|0.021249760666668692|
|         8|   1266|0.003847568221396...|
+-----+-----+-----+
```

4. Task4截图



```
In [1]: from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('random_forest').getOrCreate()
df=spark.read.csv('hdfs://localhost:9000/Task12/feature.csv',inferSchema=True,header=True)

In [2]: from pyspark.ml.feature import VectorAssembler
df_assembler = VectorAssembler(inputCols=['age_range', 'gender', 'total_logs', 'unique_item_ids', 'categories', 'browse_
df = df_assembler.transform(df)

In [3]: model_df = df.select(['features', 'label'])
train_df, test_df=model_df.randomSplit([0.75,0.25])

In [4]: from pyspark.ml.classification import RandomForestClassifier
rf_classifier=RandomForestClassifier(labelCol='label',numTrees=120).fit(train_df)
rf_predictions=rf_classifier.transform(test_df)

In [5]: from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
rf_accuracy=MulticlassClassificationEvaluator(labelCol='label',metricName='accuracy').evaluate(rf_predictions)
print('MulticlassClassificationEvaluator 随机森林测试的准确性: {0:.0%}'.format(rf_accuracy))
MulticlassClassificationEvaluator 随机森林测试的准确性: 94%

In [6]: rf_auc=BinaryClassificationEvaluator(labelCol='label').evaluate(rf_predictions)
print('BinaryClassificationEvaluator 随机森林测试的准确性: {0:.0%}'.format(rf_auc))
BinaryClassificationEvaluator 随机森林测试的准确性: 63%
```

五、问题总结

1. pyspark连接jupyter时出错

解决方案：

参考<https://blog.csdn.net/hecongqing/article/details/85016154>
<https://blog.csdn.net/donaldsy/article/details/96194346>

修改配置文件 `vim ~/.jupyter/jupyter_notebook_config.py`
找到 `#c.NotebookApp.allow_root = False`，去掉#，并修改为True

六、可能的改进之处

1. 第四题预测回头客没来得及用其他方法进行预测对比分析，后续有时间可以尝试别的预测方法